ACTUAL AUDITED TRAINING CHECKPOINT

First Quadruple training pair complete

Fixed snapshot: plan457, 2026-11-09 06:12:46 UTC. New local report; not yet uploaded.

1 of 9 goal-seed pairs complete. 2 of 18 training phases complete.
Eight BC and eight QV phases remain unrun: 800,000 target updates.

This is a **PARTIAL training and reproducibility checkpoint**. Quadruple native evaluation is **UNRUN** for BC, original QGF, all eight DFlow variants, original IPG and FK. UNRUN does not mean zero successes. No algorithm-performance or study-completion claim follows from this pair.

Exact pilot configuration

Scope	Configuration
Tasks and first lineage	Official Cube Quadruple Tasks3/4/5; first completed pair is Task3 / seed11.
Observed / action representation	55 observation dimensions; 5 native controls per step; H=5 action chunk flattened to 25 dimensions.
Training adaptation	128 x 128 networks; flow steps F=8; temperature T=1; local official 1M TRAIN shard000; 996,000 full-H5 eligible starts per goal.
Credit and paper contrast	Fresh 50,000 BC + 50,000 goal-specific QV updates; unaugmented data. This limited pilot is not the paper100M / 500k / F10 setup.

Three actual attempts, with the failed prefix preserved

Attempt	Actual outcome	Parent seconds
Fresh BC	50,000 complete; 51 checkpoint pairs; audited and reconciled.	911.232648
Original QV	FAILED / bounded partial: 15,723 returned; 15,000 durable; 723 later discarded and repeated. Stop trigger unknown.	292.626877
Exact continuation	35,000 new original-index updates; trusted15k -> effective50k; 36 checkpoint pairs; audited and reconciled.	603.704228

The continuation saved **70,446 PASS / 0 FAIL** predicates. Its 75 full-state proofs comprise **1 fresh-template zero-credit proof + 74 proofs with actor FULL TrainState and RNG frozen**. Final disk restore passed semantic and canonical-blob equality. Final payload **33bf6e496d84...**; manifest **de0c33a88fa1...**. These are reproducibility gates, not competence or Q-truth tests.

ALL REQUESTED GOALS AND SEEDS RETAINED

One pair done; eight pairs remain

No goal, seed, phase, numerical gate or primary-study arm is removed by this checkpoint.

Goal / official task	Seed	BC50k	QV50k
3 / pnp_from_square	11	COMPLETE	COMPLETE*
3 / pnp_from_square	23	UNRUN	UNRUN
3 / pnp_from_square	37	UNRUN	UNRUN
4 / cycle	11	UNRUN	UNRUN
4 / cycle	23	UNRUN	UNRUN
4 / cycle	37	UNRUN	UNRUN
5 / stack	11	UNRUN	UNRUN
5 / stack	23	UNRUN	UNRUN
5 / stack	37	UNRUN	UNRUN

* Goal3/seed11 QV reached effective50k via the audited exact15k + 35k continuation. Its combined physical QV returns are 50,723, including the known723 repeats. The original failed attempt and its unknown stop trigger remain in the record.

What the data and source checks establish

P1 covered all three goal labels over the full1M TRAIN starts. Each goal has 996,000 critic-eligible / actor-full5 starts and no native-short prefixes. All three strict joint-goal TRAIN counts are zero. This is descriptive dataset support, **not policy competence, impossibility, Q-truth or a hidden full-goal BC-success cutoff**. No BC sharing across goals is adopted.

Corrected C0 three-context source **e2d263fa / PREP cc38380e** has SCI **e737870e** and OPS **f5a00c7a** PASS, accepted SOURCE-only at Root457. It links complete BC, old failed QV15k and resumed complete QV50k. C0 is **UNRUN** at this snapshot; helper, fresh admission and any actual engineering execution are separate. The old blocked source and review history remain immutable.

Remaining evaluation dependencies

After the remaining training and exact frozen-parent checks, each family still needs its own admitted callback / H5 boundary / numerical readiness checks, validation and fresh confirmation. Original QGF and FK can follow their own reviewed gates; DFlow and IPG do not inherit a numerical waiver. No native policy-success result is supplied here.

The primary Task2 study still requires **all33 guided conditions** before main100: 30 are frozen and IPG3 remain unqualified. Fresh confirmation and matched held-out evaluation remain distinct. New main100 held-out episodes = **0**; the 540 matched curve points x100 remain pending. Quad training is not a substitute for that barrier.

Fresh40 outcomes, own selected settings

Successes / 40 on seeds11,23,37. Different selected settings; this table is not held-out or a causal averaging comparison.

Condition	lambda	seed11	seed23	seed37
Regular / raw Q	0	35/40	40/40	37/40
Regular / raw Q	1	28/40	40/40	40/40
Stochastic / raw Q	0	40/40	40/40	39/40
Stochastic / raw Q	1	37/40	40/40	40/40
Regular / averaged Q	0	38/40	40/40	37/40
Regular / averaged Q	1	25/40	40/40	40/40
Stochastic / averaged Q	0	40/40	40/40	40/40
Stochastic / averaged Q	1	35/40	40/40	40/40
Original QGF	-	40/40	40/40	40/40
Original FK	-	40/40	40/40	40/40
Original IPG	-	Unqualified	Unqualified	Unqualified

QGF and FK are original families. DFlow rows retain regular/stochastic x raw/averaged-Q x lambda0/1 identities. The selected kappa, eta, iteration count and related settings differ across rows and seeds. Therefore these scores do **not establish that averaging helps, a tuned DFlow advantage, Q-truth, or a held-out win**. IPG unqualified is not zero success.

Full tables and cohort scopes remain available

The unchanged plan365 tables retain all24 selected DF tuples [D1], all72 literal fresh40 DF alternatives [D4], QGF/FK selected tuples [D2], full27,720 DF TRAIN score rows [D5], QGF648 Task2/Task5 rows [D6] and FK2,160 TRAIN rows [D7]. The DF TRAIN scores are repeated scores on one paired TRAIN16 panel, not native episodes.

The newer saved QGF Task2 sweep view [Q1-Q2] covers **324 validation configurations / 6,480 episode views**: 1,620 inherited M8 reference episodes + 4,860 new joint-scan episodes. Nine fresh40 confirmation alternatives represent360 episodes. Inherited and new views are explicitly separated; no new queries were made to build the table view.

Full selected configuration fields and immutable CSV paths are in **SOURCE_PINS.json** and the copied small tables under **evidence/**. This report does not replace, overwrite or reselect those old tuples. Task5 QGF confirmation40/36/40, old DF numerical failures and IPG failures remain unchanged.

Measured time and preserved evidence

All clocks below are fixed-plan457 saved accounting, not a promise that the remaining study fits.

Accounting category	Seconds
Measured completed scientific outer-parent, once	48,775.114103380125
Separate UNKNOWN conservative charge; not measured time	600
Used / committed; active scientific caps = 0	49,375.114103380125
Remaining from initial86,400-second tranche	37,024.885896619875

BC911.232648 + old QV292.626877 + resumed QV603.704228 seconds are already included once. Worker, watchdog, update-body, setup, helper, acquisition, archive and publication/admin clocks are nonadditive. Future600/7200 caps and disk/RSS/output fit require separate current admission; no bulk cap, parallel speedup or new phase is authorized by this PDF.

Evidence key; full SHA256 and exact paths are in SOURCE_PINS.json

Saved evidence	SHA prefixes / scope
BC audit + required scope / Root442	d82d4bf4 + 7d3e97c0 / c9576084
Failed QV audit + required scope / Root447	406bc922 + 27fc68db / 013b525a
Resume audit + mandatory scope and addendum / Root456	fffb8d2c + c010520a + 246d4b42 / 5ca3c57a
C0 schema compatibility companion; no new checks	c30a2374 + a5af163d
Corrected C0 source + PREP / SCI + OPS / Root457	e2d263fa + cc38380e / e737870e + f5a00c7a / 72843da9

Public archive boundaries and this new local report

The public complete-campaigns snapshot stops at **2026-11-08 23:35:08.740055 UTC / plan409**; it excludes later Quad work. The separate public 4MB P1 code + saved QA metadata release excludes BC/QV weights and raw binary arrays. The existing public PDF directory does not yet include this checkpoint.

This new PDF/TXT is not yet uploaded. The local manifest is a reference bundle, not a whole portable backup. All three audited canonical states are indexed but no checkpoint bytes, raw data, scientific NPZ or runtime are embedded; no ZIP was created. Full attempts remain external and uncut. Root independent content and all-four-page review is pending author handoff. This remains PARTIAL.